

CH6. Process Scheduling

www.ajou.ac.kr

Department of Software

Kyunghee Choi



아주대학교

Contents

-
- 1 Process Scheduling Criteria
 - 2 Basic Concepts
 - 3 Scheduling Algorithms
 - 4 Thread Scheduling
 - 5 Multiple-Processor Scheduling
 - 6 Real-time Scheduling

Scheduling/Schedule

The image shows a Google Calendar interface for the week of May 3 to May 9, 2020. The calendar grid has days from Monday to Sunday. The days from May 3 to May 6 are in light blue boxes, while May 7 is highlighted in dark blue with a white number '7'. The days from May 8 to May 9 are in light blue boxes. The hours on the left are listed as GMT+09, from 오전 10시 to 오후 5시.

- May 3: No events
- May 4: No events
- May 5: 어린이날 (Mother's Day) - (제목 없음) 오전 10시~오후 2시
- May 6: No events
- May 7: My Schedule (highlighted)
- May 8: 어버이날 (Father's Day)
- May 9: No events

On the right side of the calendar, there are several icons: a magnifying glass, a question mark, a gear, a dropdown menu labeled '주' (Week), a grid icon, the 'CONNECTING MINDS' logo, and a blue circular button labeled '경회' (Meeting).

Google Calendar

Scheduling/Schedule

< 진도 계획 >

주	강의 주제	언어	담당교수	수업방법	평가방법
1	강의소개 및 운영체제 개요	한	최경희	강의	
2	운영체제 구조	한	최경희	강의	
3	프로세스 관리 I	한	최경희	강의	
4	프로세스 관리 II	한	최경희	강의	
5	쓰레드	한	최경희	강의	
6	프로세스 동기화 I				
7	프로세스 동기화 II				
8	증간고사				
9	CPU 스케줄링				
10	데드락 관리				

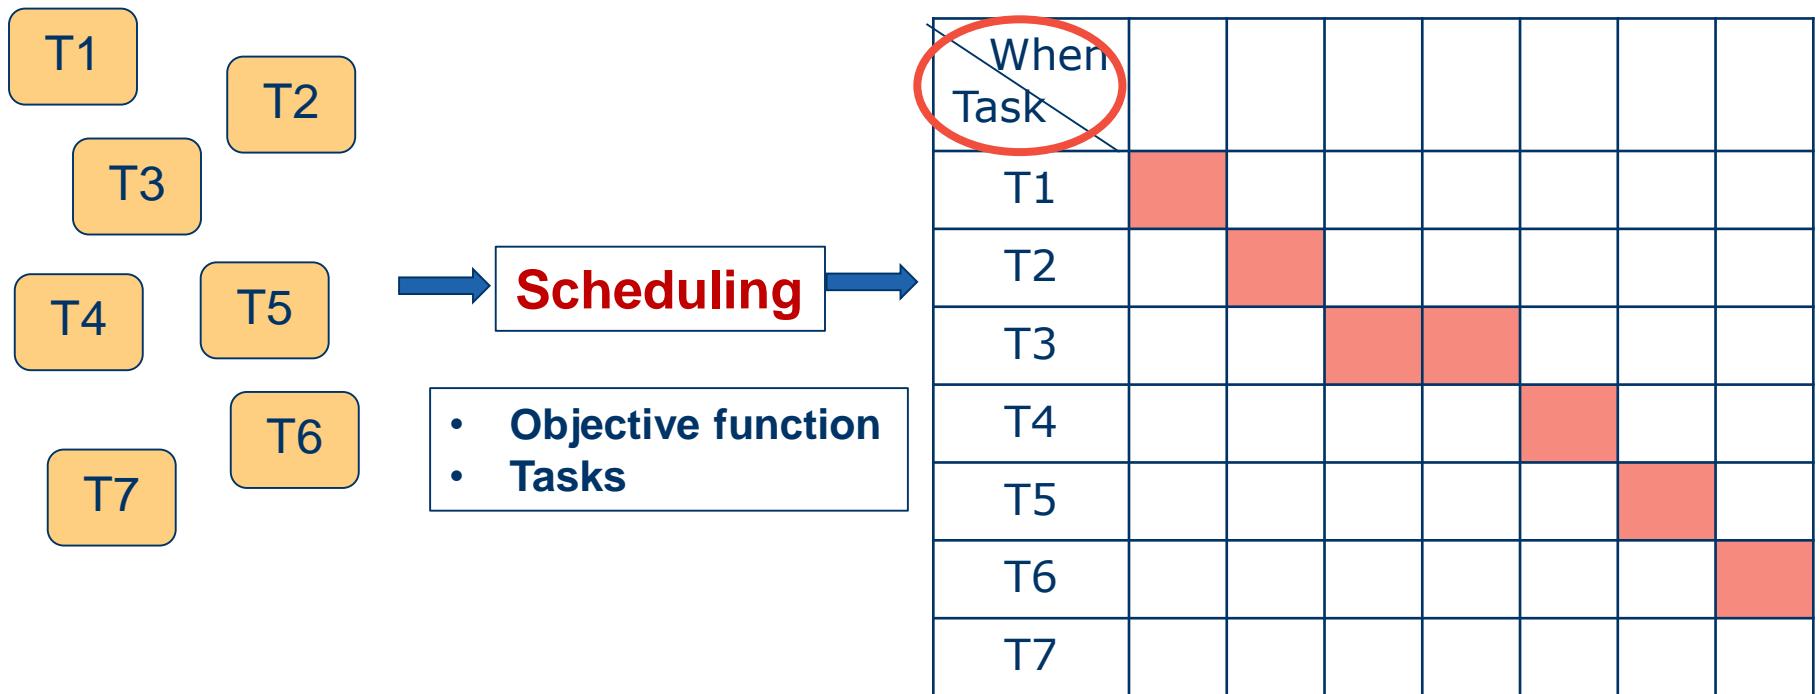
운영체제 강의 스케줄

Scheduling/Schedule

- ❖ 작업 스케줄링
- ❖ 공정스케줄링
- ❖ **Scheduling in Kernel**
 - Packet Scheduling
 - IO Scheduling
 - CPU scheduling - process scheduling
 - Job Scheduling
- ❖ **Scheduling in Multimedia Application System**

Scheduling/Schedule

Key Points



1

Process Scheduling Criteria

Process Scheduling Criteria

❖ CPU utilization

- keep the CPU as busy as possible (0-100%)

❖ Throughput

- the number of processes that complete their execution per time unit

❖ Turnaround time

- amount of time to execute a particular process
(Turnaround time = **completion** time - **submission** time)

❖ Waiting time

- amount of time a process has been waiting **in the ready queue**
(neither CPU execution nor I/O blocking)

❖ Response time

- amount of time it takes from when a **request** was submitted until the **first response** is produced, **not output** (for time-sharing environment)

Optimization Criteria

❖ Optimization Criteria

- Maximize CPU utilization and throughput
- Minimize turnaround time, waiting time, and response time

❖ In most cases,

- optimize the average measures

❖ Under some circumstances,

- optimize the minimum or maximum values, rather than average

❖ For interactive systems

- minimize the variance in the response time
→ a system with reasonable and predictable response time

2

Basic Concepts

Basic Concepts

❖ Objective of multiprogramming

- Maximize CPU utilization

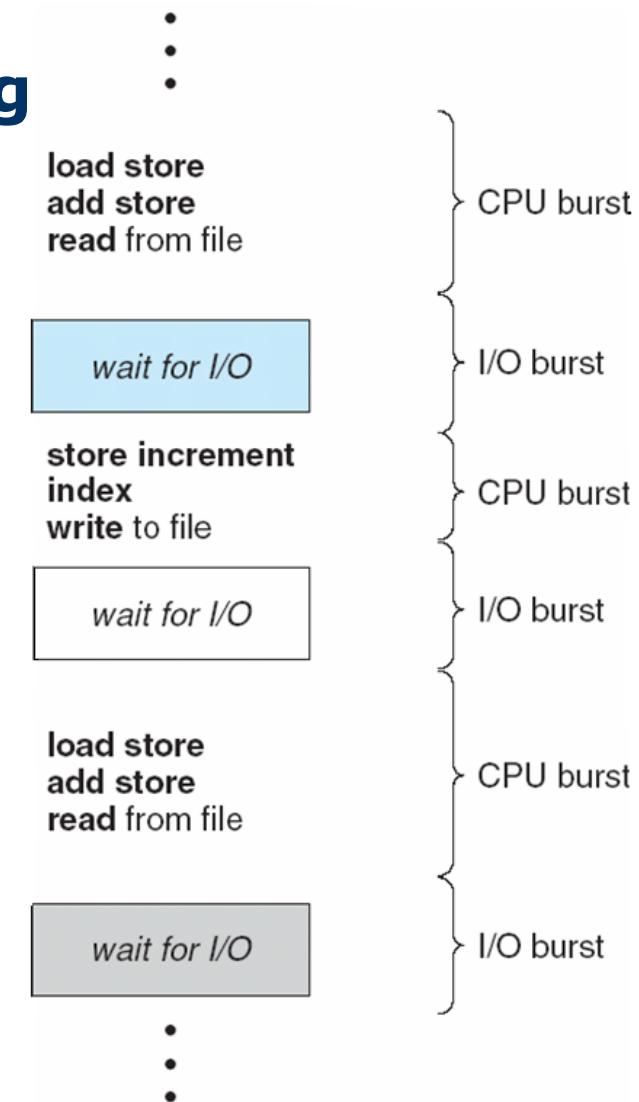
❖ CPU-I/O Burst Cycle

- Process execution consists of a cycle of CPU execution and I/O wait

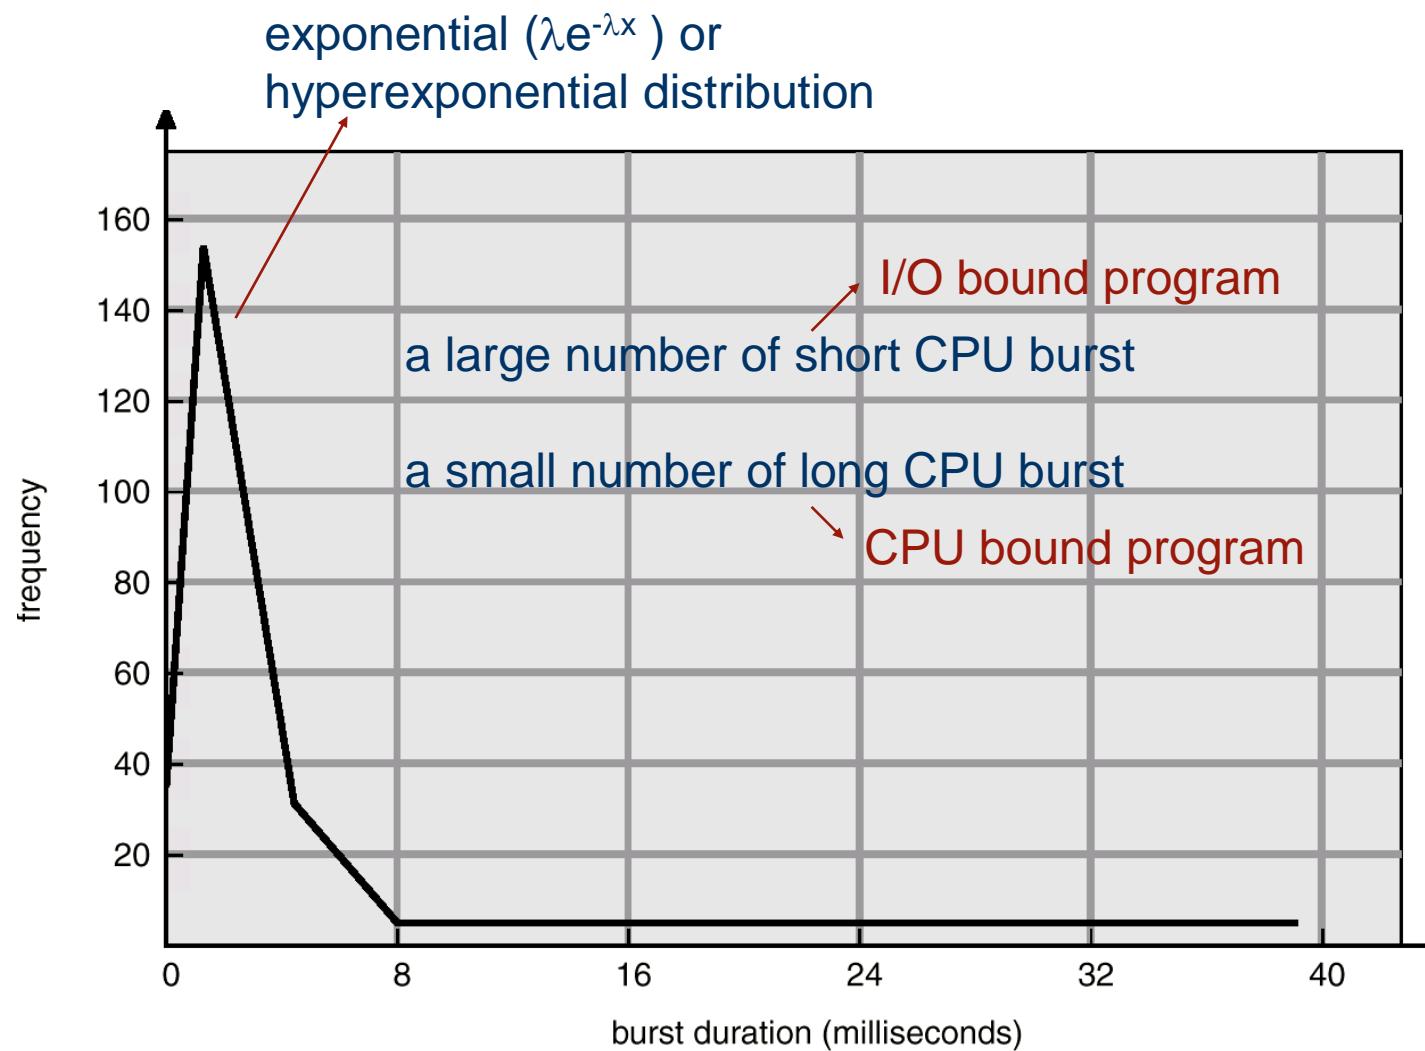


❖ CPU burst distribution

- (see next page)



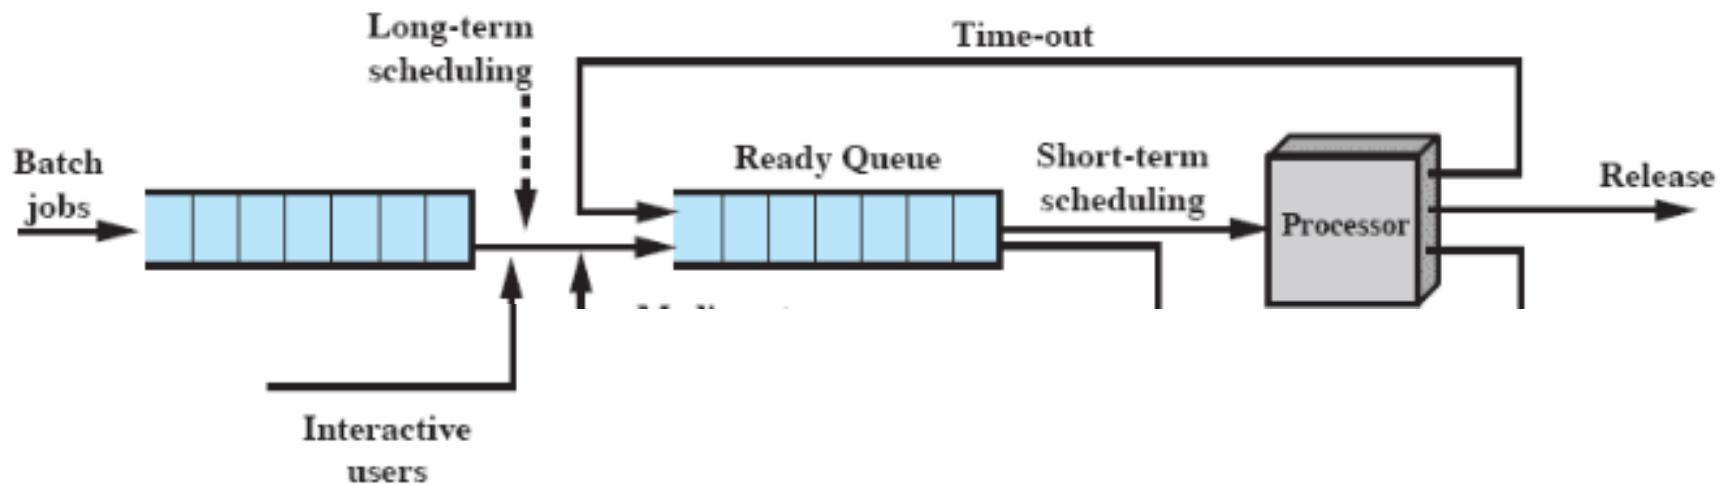
Histogram of CPU-burst Times



CPU Scheduler

❖ CPU scheduler (short-term scheduler)

- Selects one of the processes in the ready queue, and allocates the CPU to that process.



When CPU Scheduling may take place ?

❖ CPU scheduling decisions may take place when a process:

1. Switches from **running** to **waiting** state
(e.g. I/O wait, child termination wait)
 2. Switches from **running** to **ready** state (e.g. time-out)
 3. Switches from **waiting** to **ready** state (e.g. I/O completion)
 4. **Terminates**
-
- case 1, 4 => there is **no choice** in terms of scheduling.
a new process must be selected
 - case 2, 3 => there is a **choice**.
(a current process is in ready state)

Preemptive and Nonpreemptive Scheduling

❖ **Nonpreemptive (cooperative) scheduling:**

- scheduling takes place *only* under cases 1 and 4
 - A process in the running state will continue until it terminates or blocks itself
- (ex) windows 3.x, old Mac OS
- It is the only method that can be used on hardware platforms which do not have the special hardware (e.g. timer)

❖ **Preemptive scheduling:**

- scheduling may take place under all cases
- (ex) Most operating systems
- It requires a special hardware (e.g. timer)

Problems of Preemptive scheduling

- ❖ **Preemptive scheduling incurs a cost associated with coordination of access to shared data → synchronization mechanism (Chap 5)**
- ❖ **Race Condition**
 - When data are shared, one process updates that data and is preempted. After that, the second process tries to read that data
- ❖ **Preemption in user mode**
 - (e.g.) Two processes share data
 - While one is updating data, it is preempted and another tries to read or modify the data → in an inconsistent state
- ❖ **Preemption in kernel mode:**
 - (e.g.) All kernel routines share the kernel data
 - The kernel may process a system call on behalf of a process and involves changing important kernel data.
The process is preempted during these changes and the kernel (or device driver) to read or modify the same data.
→ in an inconsistent state (more dangerous)

Problems of Preemptive scheduling

❖ Certain OSs simply deal with the preemption problem in kernel

- by waiting for
 - (1) a system call to complete, or
 - (2) a I/O block to take place before doing context-switching
- This kernel-execution model is poor for supporting real-time computing and multi-processing.

Dispatcher

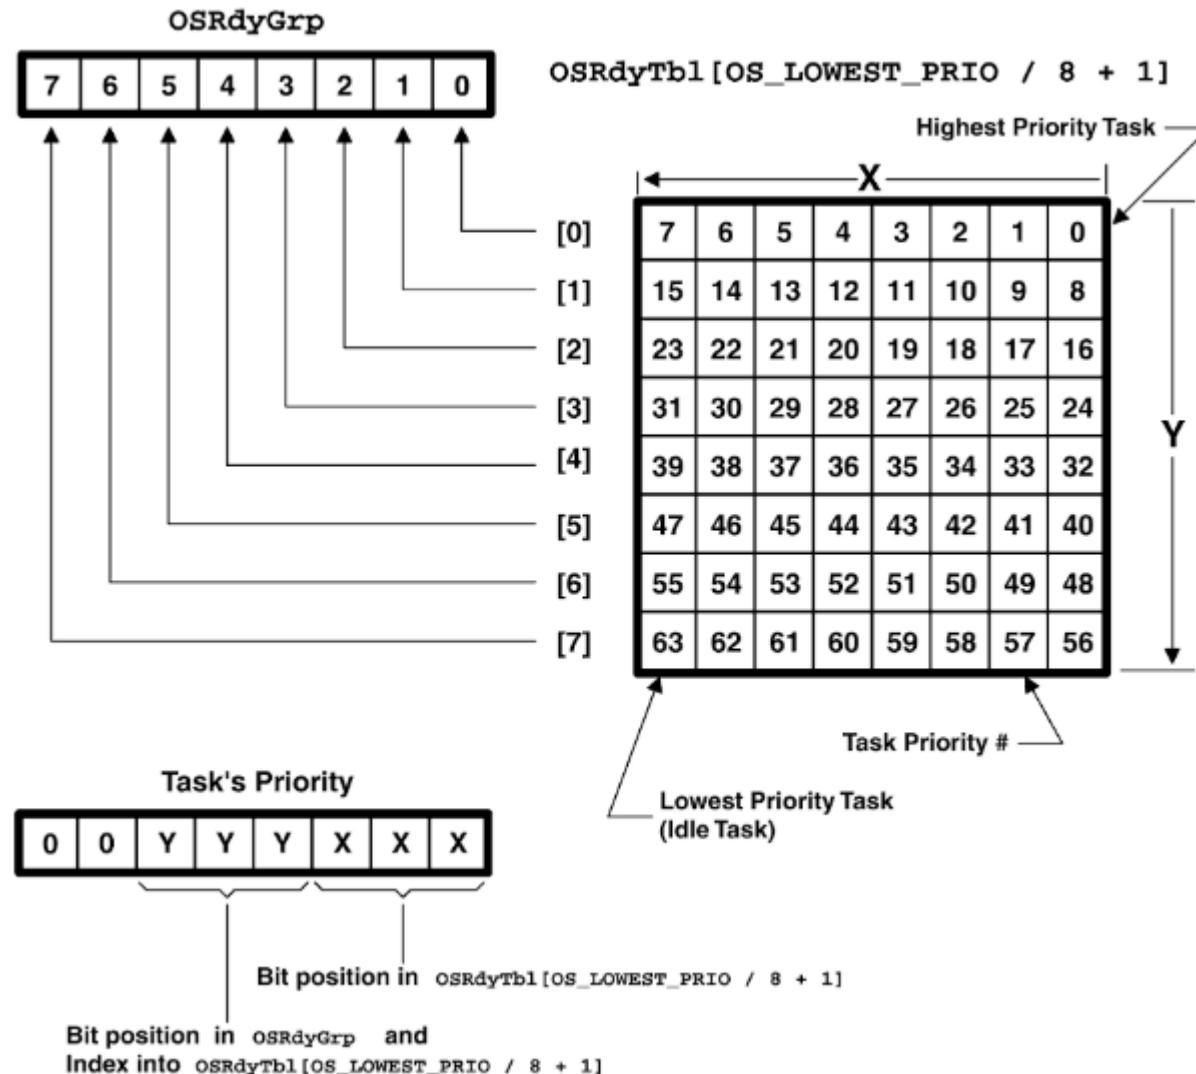
❖ Dispatcher

- gives control of the CPU to the process selected by the short-term scheduler
- this function involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program

❖ Dispatch latency

- the time it takes for the dispatcher to *stop* one process and *start* another running.
- dispatch latency should be fast as possible

Dispatcher Example(uC/OS)



Dispatcher Example

```
OS_ENTER_CRITICAL();
if ((OSIntNesting == 0) && (OSLockNesting == 0)) {
    y           = OSUnMapTb1[OSRdyGrp];
    OSPrioHighRdy = (INT8U)((y << 3) + OSUnMapTb1[OSRdyTb1[y]]);
    if (OSPrioHighRdy != OSPrioCur) {
        OSTCBHighRdy = OSTCBPrioTb1[OSPrioHighRdy];
        OSCtxSwCtr++;
        OS_TASK_SW();
    }
}
OS_EXIT_CRITICAL();
```



3

Scheduling Algorithms

Scheduling Algorithm

- ❖ **First-Come First-Serve(FCFS) scheduling**
- ❖ **Shortest-Job-First (SJR) Scheduling**
- ❖ **Priority Scheduling**
- ❖ **Round Robin (RR)**
- ❖ **Multilevel Queue Scheduling**

First-Come, First-Served (FCFS) Scheduling

❖ Example:

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

❖ Arrival order: P_1, P_2, P_3 (arrival time t=0)

- The *Gantt Chart* for the schedule is:

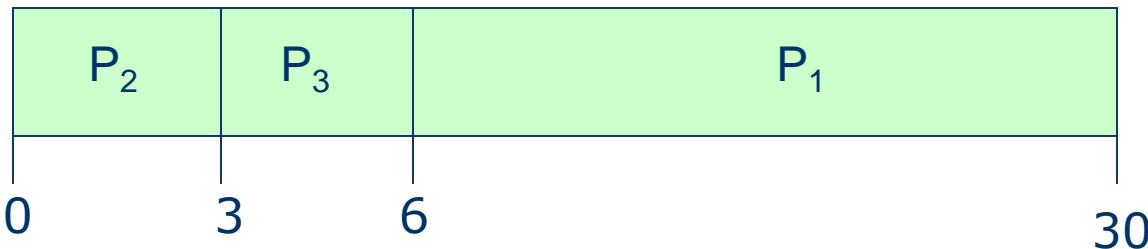


- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)

❖ Arrival order: P_2, P_3, P_1 . (arrival time t=0)

The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$, $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
→ Much better than previous case.

❖ Convoy effect

- short processes behind long process →
short processes wait for the one long process to get off the CPU
- This effect results in lower CPU and device utilization

Shortest-Job-First (SJF) Scheduling

❖ Shortest-Job-First (SJF) scheduling

(**Shortest Process Next: SPN, Shortest Request Next: SRN**)

- each process is associated with the length of its next CPU burst
- When the CPU is available,
schedule the process with the *shortest* next CPU burst
- it is difficult to know the next CPU burst → (solution) *prediction*

❖ Two schemes:

- Nonpreemptive – The currently running process *cannot be preempted* until completes its CPU burst
- Preemptive –
if CPU burst of a new process < remaining time of a current process,
then preempt → **Shortest-Remaining-Time-First (SRTF)**
“preemptive SJF = SRTF”

❖ SJF is optimal

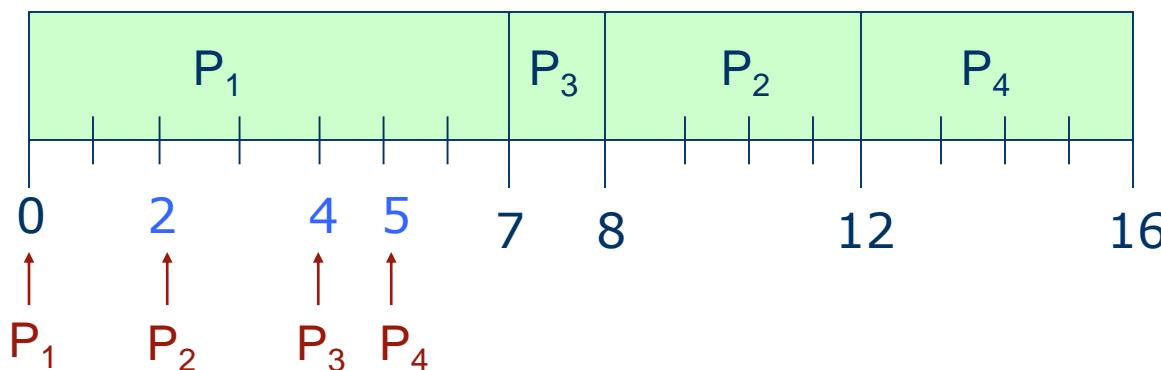
- gives minimum average waiting time for a given set of processes

Example of Non-Preemptive SJF

❖ Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

❖ SJF (non-preemptive)



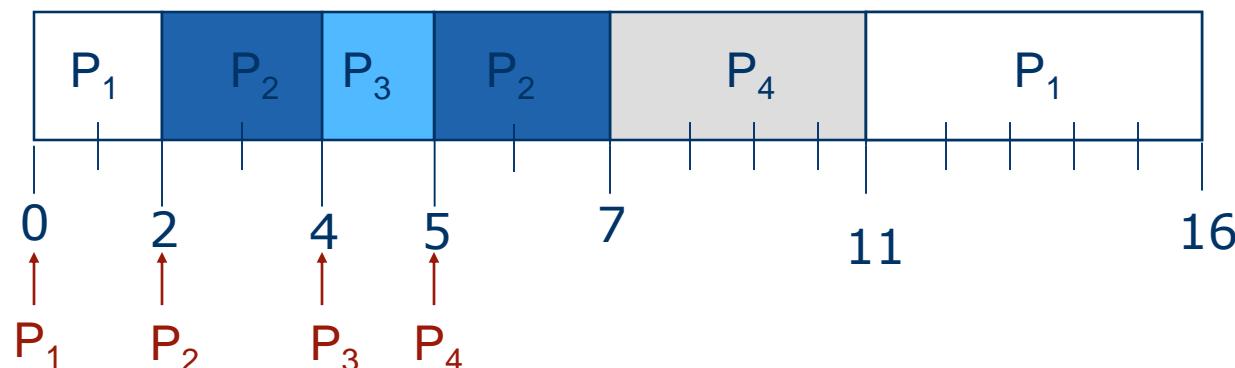
- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

❖ Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

❖ SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Priority Scheduling

❖ Priority Scheduling

- Each process is associated with a priority number (integer)
- The CPU is allocated to the process with the **highest** priority (usually, smallest priority number \equiv highest priority)

❖ Priorities can be defined

- internally: time limits, memory requirements, # of open files, ratio of I/O to CPU
- externally: importance of process, fund type and amount, political factor

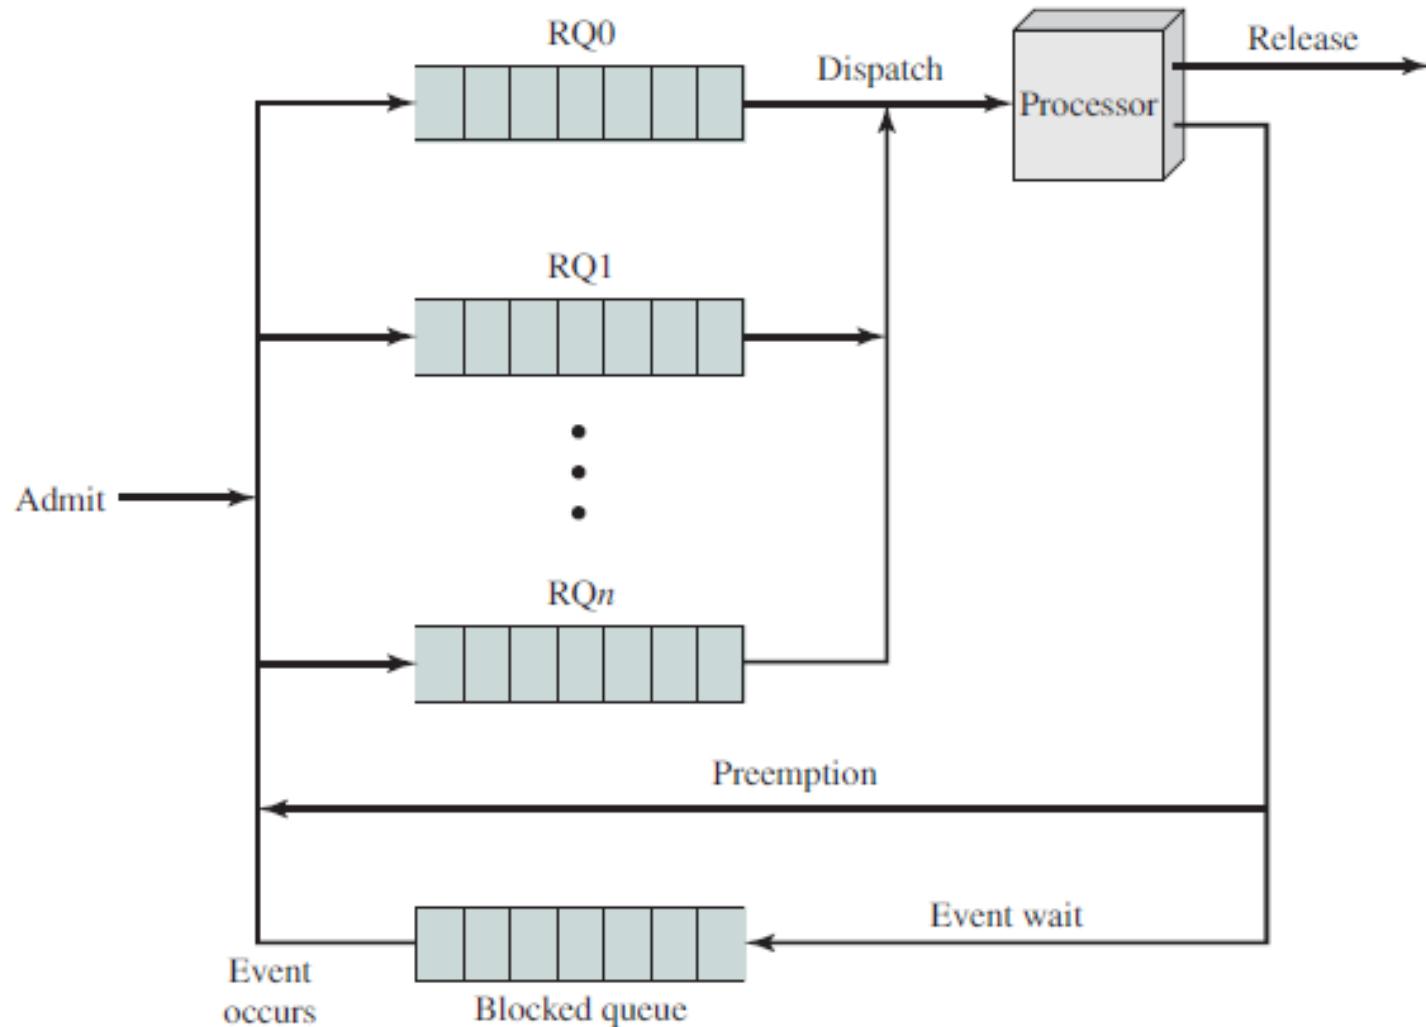
❖ Two schemes

- preemptive
- nonpreemptive

❖ SJF is a special case of priority scheduling

- priority = the predicted next CPU burst time

Priority Queueing



Starvation

❖ Problem of priority scheduling → Starvation (Indefinite blocking)

- low priority processes may *never* execute.

❖ Solution of starvation → aging

- Aging: a technique of gradually increasing the priority of processes that wait in the system for a long time

Round Robin (RR) Scheduling

❖ Round Robin scheduling (processor sharing)

- Each process gets a small unit of CPU time (time quantum), usually 10-100 milliseconds
- After this time has elapsed, the process is preempted and added to the end of the ready queue

❖ Maximum waiting time

- For n ready processes, time quantum q ,
max. waiting time = $(n-1)q$

❖ Performance

- q large \Rightarrow FIFO (FCFS)
- q small $\Rightarrow q$ must be large with respect to context switch time, otherwise, overhead is too high

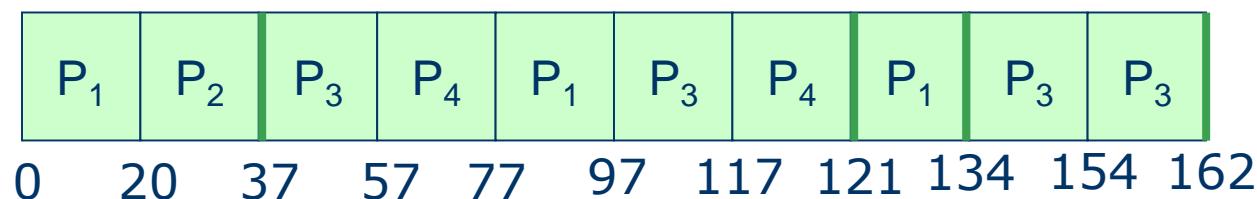


Example: RR with Time Quantum = 20

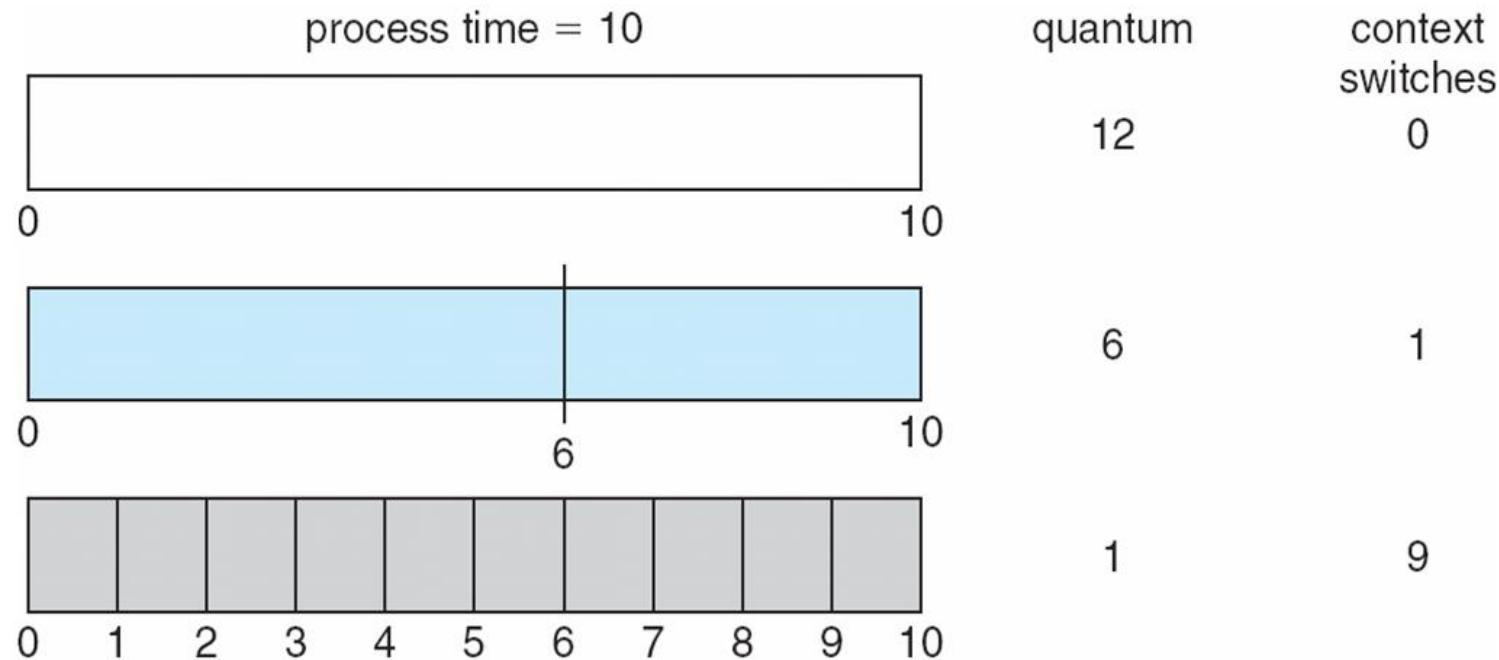
❖ Example

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

❖ The Gantt chart is:

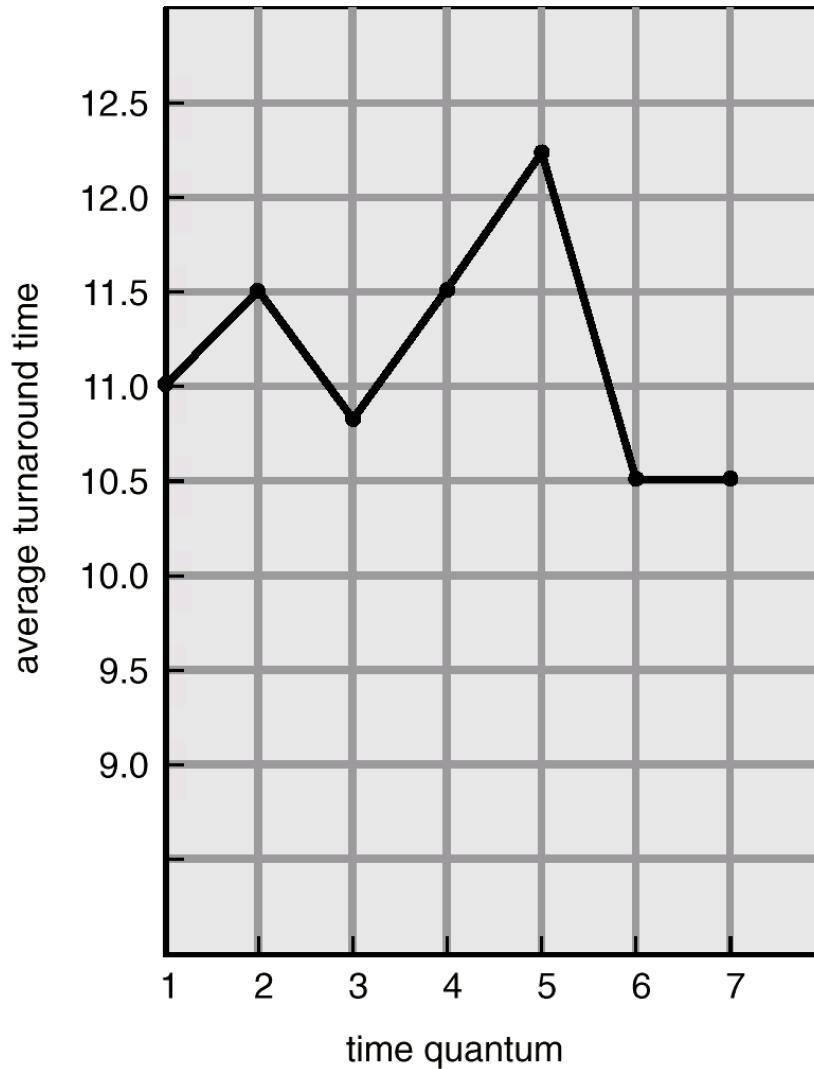


Time Quantum and Context Switch Time



- ❖ **Smaller time quantum increases context switches**
- ❖ **A rule of thumb:**
 - 80% of CPU bursts should be shorter than quantum time

Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

- ❖ **turnaround time does not necessarily improve as time-quantum size increases**

Multilevel Queue Scheduling

❖ Multilevel Queue Scheduling

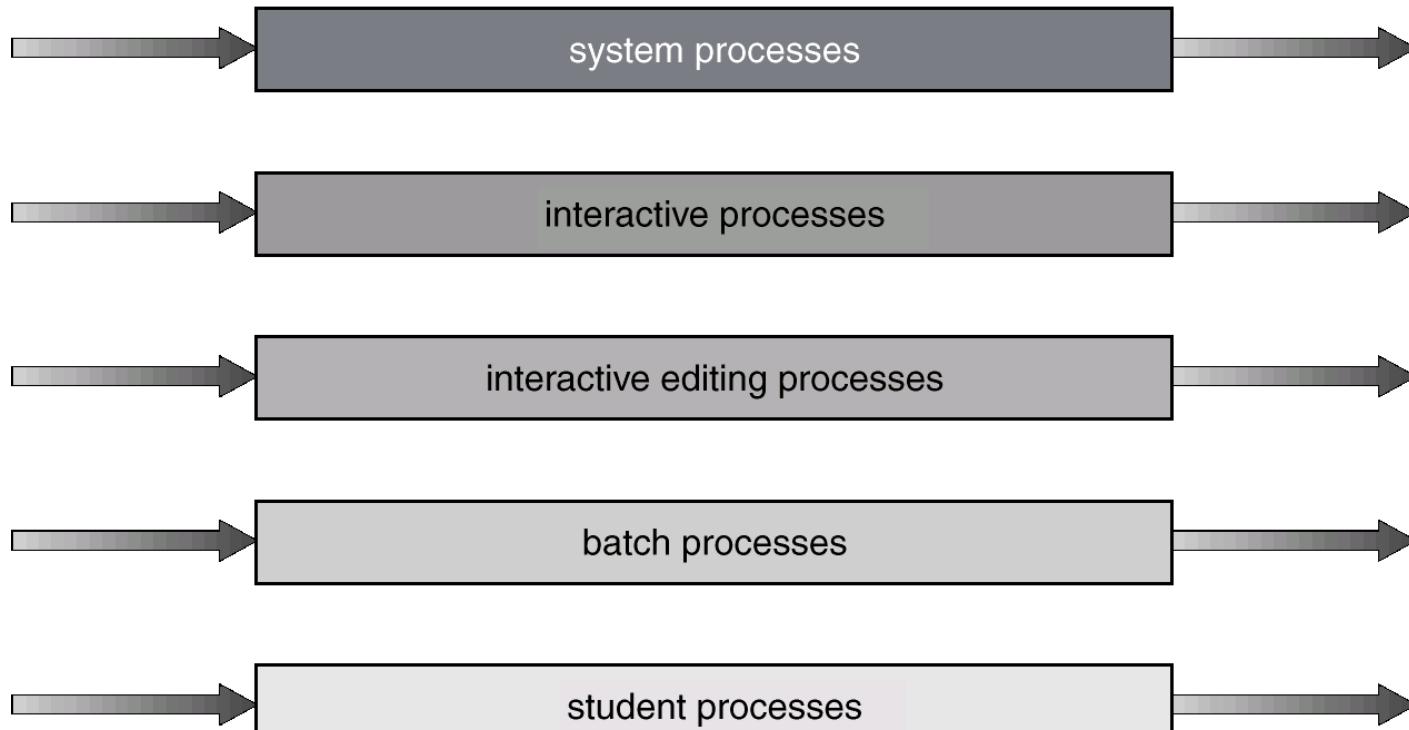
- Ready queue is partitioned into separate queues:
- Each queue has its *own* scheduling algorithm,
e.g. foreground (interactive) – RR
background (batch) – FCFS

❖ Scheduling must be done between the queues

- Fixed priority scheduling
 - serve all from foreground, then from background
 - Possibility of *starvation*
- Time slice
 - each queue gets a certain amount of CPU time which it can schedule amongst its processes
 - e.g. 80% to foreground in RR, 20% to background in FCFS

Multilevel Queue Scheduling

highest priority



lowest priority

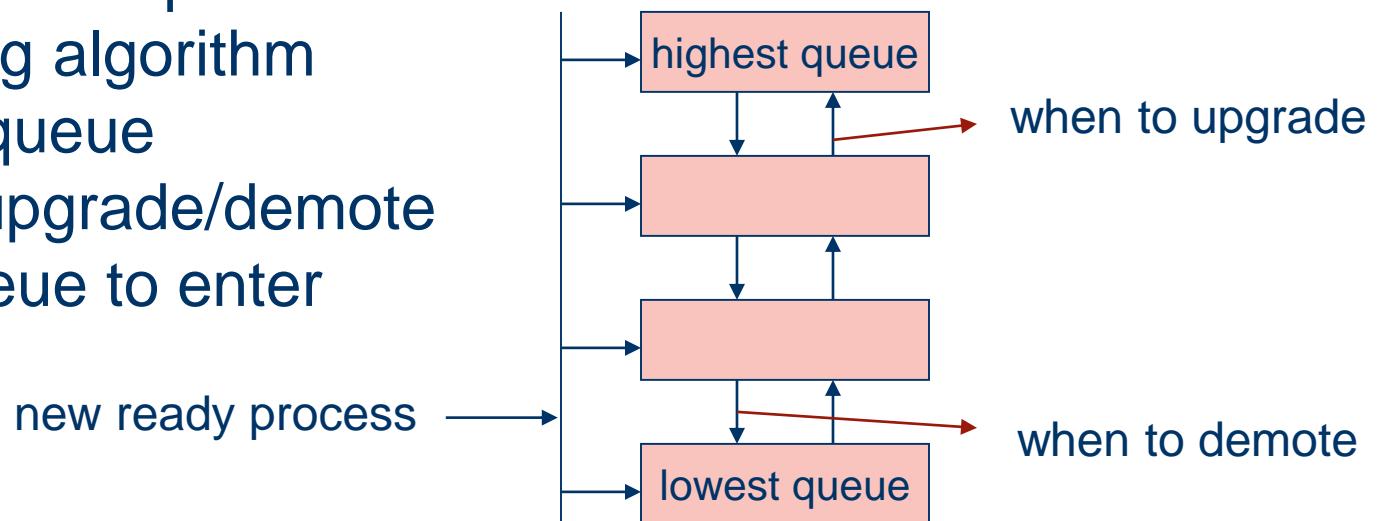
Multilevel Feedback Queue Scheduling

- ❖ **A process can move between the various queues**

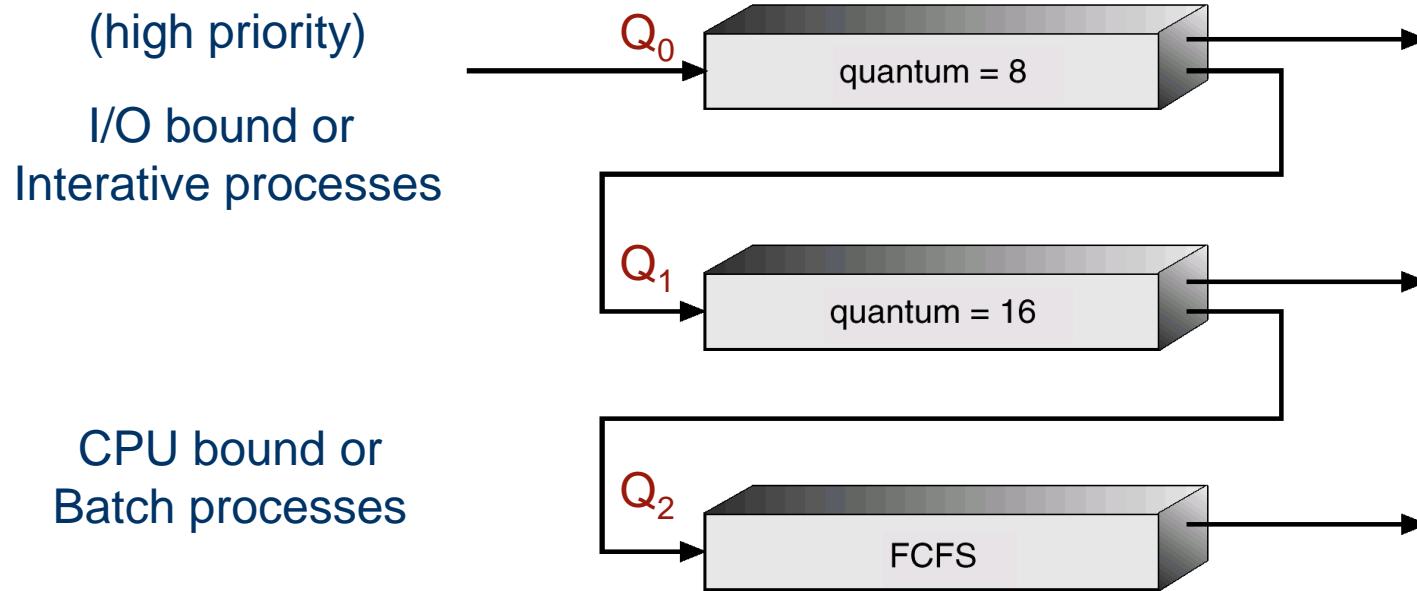
- Aging can be implemented this way → flexible

- ❖ **Multilevel-feedback-queue scheduler defined by the following parameters:**

- the number of queue
- scheduling algorithm for each queue
- when to upgrade/demote
- which queue to enter



Multilevel Feedback Queues



❖ Example:

- A new job => queue Q_0 ($q=8$, FCFS)
- If it does not finish in 8 ms, the job => queue Q_1 ($q=16$, FCFS)
- If it still does not complete, the job => queue Q_2
- priority: $Q_0 > Q_1 > Q_2$
 - The process in Q_2 are run, only when Q_0 and Q_1 are empty



4

Thread Scheduling

Thread Scheduling

- ❖ **Distinction between user-level threads and kernel-level threads**
- ❖ **Process-contention scope (PCS) - for user-level threads**
 - On system using many-to-one mapping or many-to-many mapping, the thread library schedules user-level threads to run on an available LWP (lightweight process)
 - competition for CPU takes place among threads in the same process
 - Typically, PCS is done according to priority
- ❖ **System-contention scope (SCS) - for kernel-level threads**
 - the kernel schedules kernel-level threads onto available CPU
 - competition for CPU takes place among all threads in the system
 - Systems using one-to-one mapping schedule threads using only SCS
- ❖ **Pthread scheduling policies - allow specifying during thread creation**
 - PTHREAD_SCOPE_PROCESS: PCS scheduling (many-to-many)
 - PTHREAD_SCOPE_SYSTEM: SCS scheduling (one-to-one)

Pthread scheduling

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
void *runner(void *param); /* the thread runs in this function */

int main(int argc, char *argv[])
{
    int i, scope;
    pthread_t tid[NUM_THREADS]; /* the thread identifier */
    pthread_attr_t attr; /* set of attributes for the thread */

    /* get the default attributes */
    pthread_attr_init(&attr);
    /* set the scheduling algorithm PROCESS or SYSTEM */
    /* On Linux, Mac OS X, only PTHREAD_SCOPE_SYSTEM is supported.
     * Solaris supports both. */
    if (pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM) != 0)
        fprintf(stderr, "Unable to set scheduling scope\n");
    /* now create the threads */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
    /* Now join on each thread */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_join(tid[i], NULL);
}

void *runner(void *param)
{
    /* do some work */
    ...
    pthread_exit(0);
}
```



5

Multiple-Processor Scheduling

Multiple-Processor Scheduling

❖ **Multiple-processor scheduling**

- CPU scheduling more complex when multiple CPUs are available

❖ **Multiprocessor system**

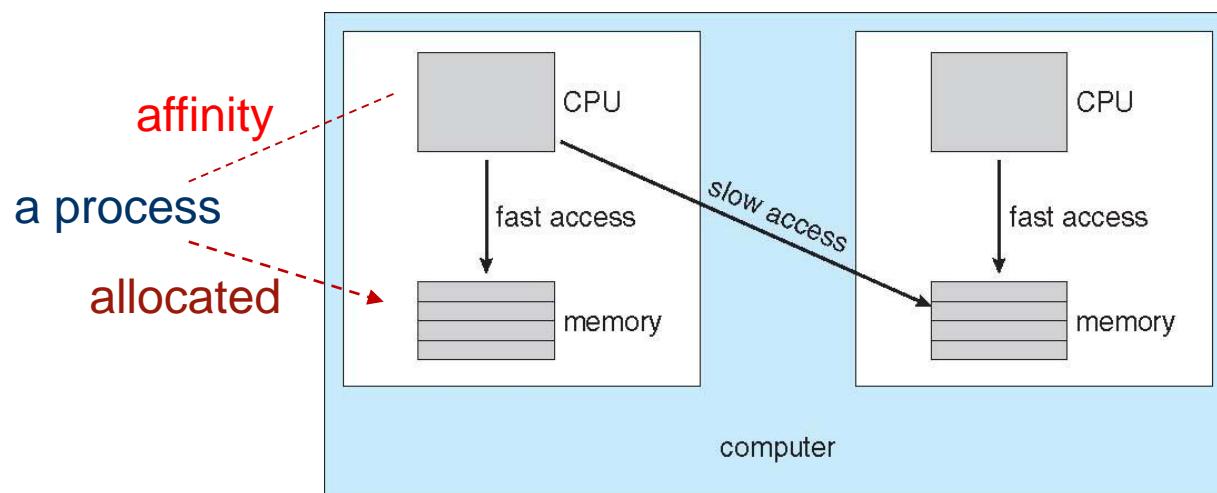
- system in which the processors are identical (homogeneous)

❖ **Multiple-processor scheduling approaches**

- **Asymmetric** multiprocessing
 - only one processor(master processor) accesses the system data structures, alleviating the need for data sharing
 - the other processors execute only user code
 - reduce the need for data sharing → simple
- **Symmetric** multiprocessing (SMP) - self-scheduling
 - each processor is self-scheduling
 - ready queue
 - a common ready queue, or
 - separate queue for each processor (in most contemporary OS's supporting SMP)

Processor Affinity

- ❖ **Processor Affinity - a process has an affinity for the processor on which it is currently running;**
 - The most SMP systems avoid migration of processes from one processor to another, instead attempt to keep a process running on the same processor
 - Because of the high cost of invalidating and repopulating caches
- ❖ **Forms of processor affinity**
 - soft affinity - possible to migrate between processors
 - hard affinity - specify a processor or processor sets
- ❖ **NUMA(non-uniform memory access) and CPU scheduling**



Load Balancing

❖ Load balancing

- keep the workload evenly distributed across all processors

❖ On systems with common ready queue

- load balancing is often unnecessary

❖ On systems with separate private queue – most contemporary OS

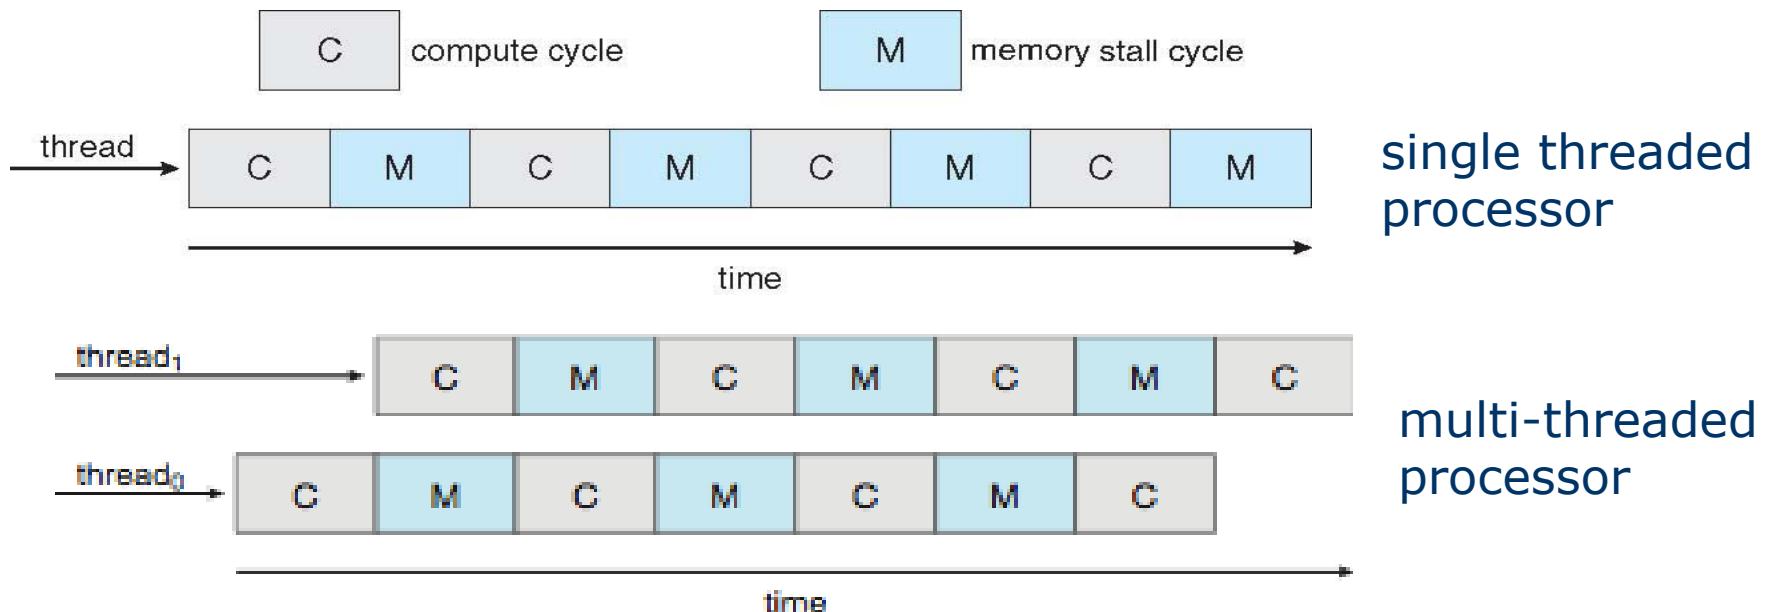
- two approaches to load balancing
 - push migration:
 - periodically checks & if an imbalance is found, then migrate processes from overloaded to idle or less-busy processors
 - pull migration:
 - an idle processor pulls a waiting task from a busy processor
- two migration approaches are not mutually exclusive
(both cannot happen at the same time)

❖ load balancing often counteracts the benefits of processor affinity

Multicore Processors

❖ Multithreaded processor

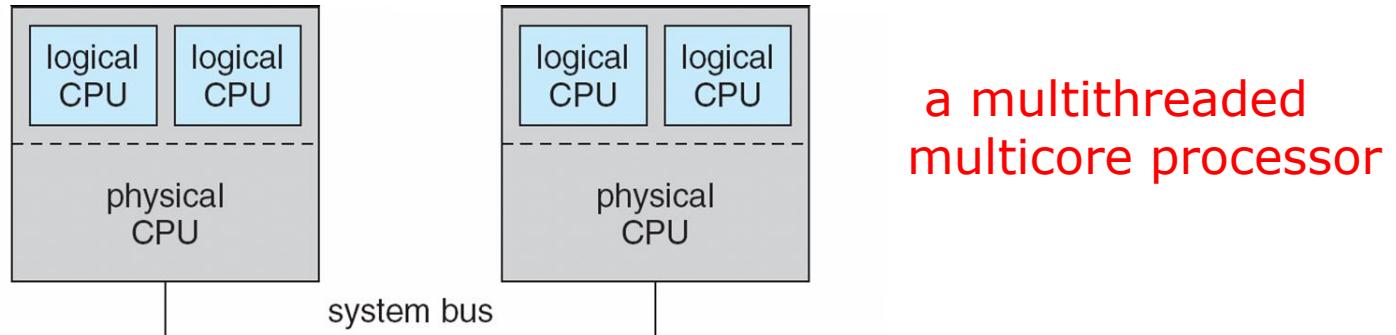
- multiple logical processor on a physical processor (core)
- each logical processor has its own registers → fast context switching
- Takes advantage of memory stall to make progress on another thread while memory retrieve happens



Multicore Processors (cont')

❖ Multicore processor

- multiple processor core on same physical chip



❖ Two different levels of scheduling

- choose which software thread to run on each hardware thread (logical processor) – by OS (use any scheduling algorithm)
 - coarse-grained multithreading
 - the cost of switching between threads is high
- choose which hardware thread to run on the core – by hardware
 - fine-grained multithreading
 - the cost of switching between threads is small



6

Real-time Scheduling

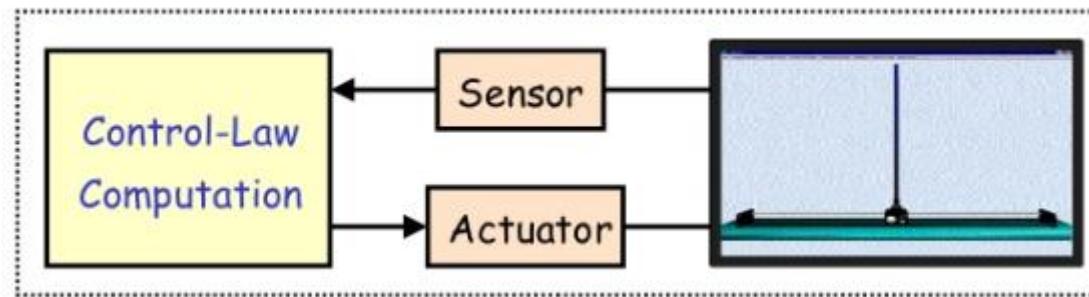
Real-time Systems

❖ Example of Real-time System

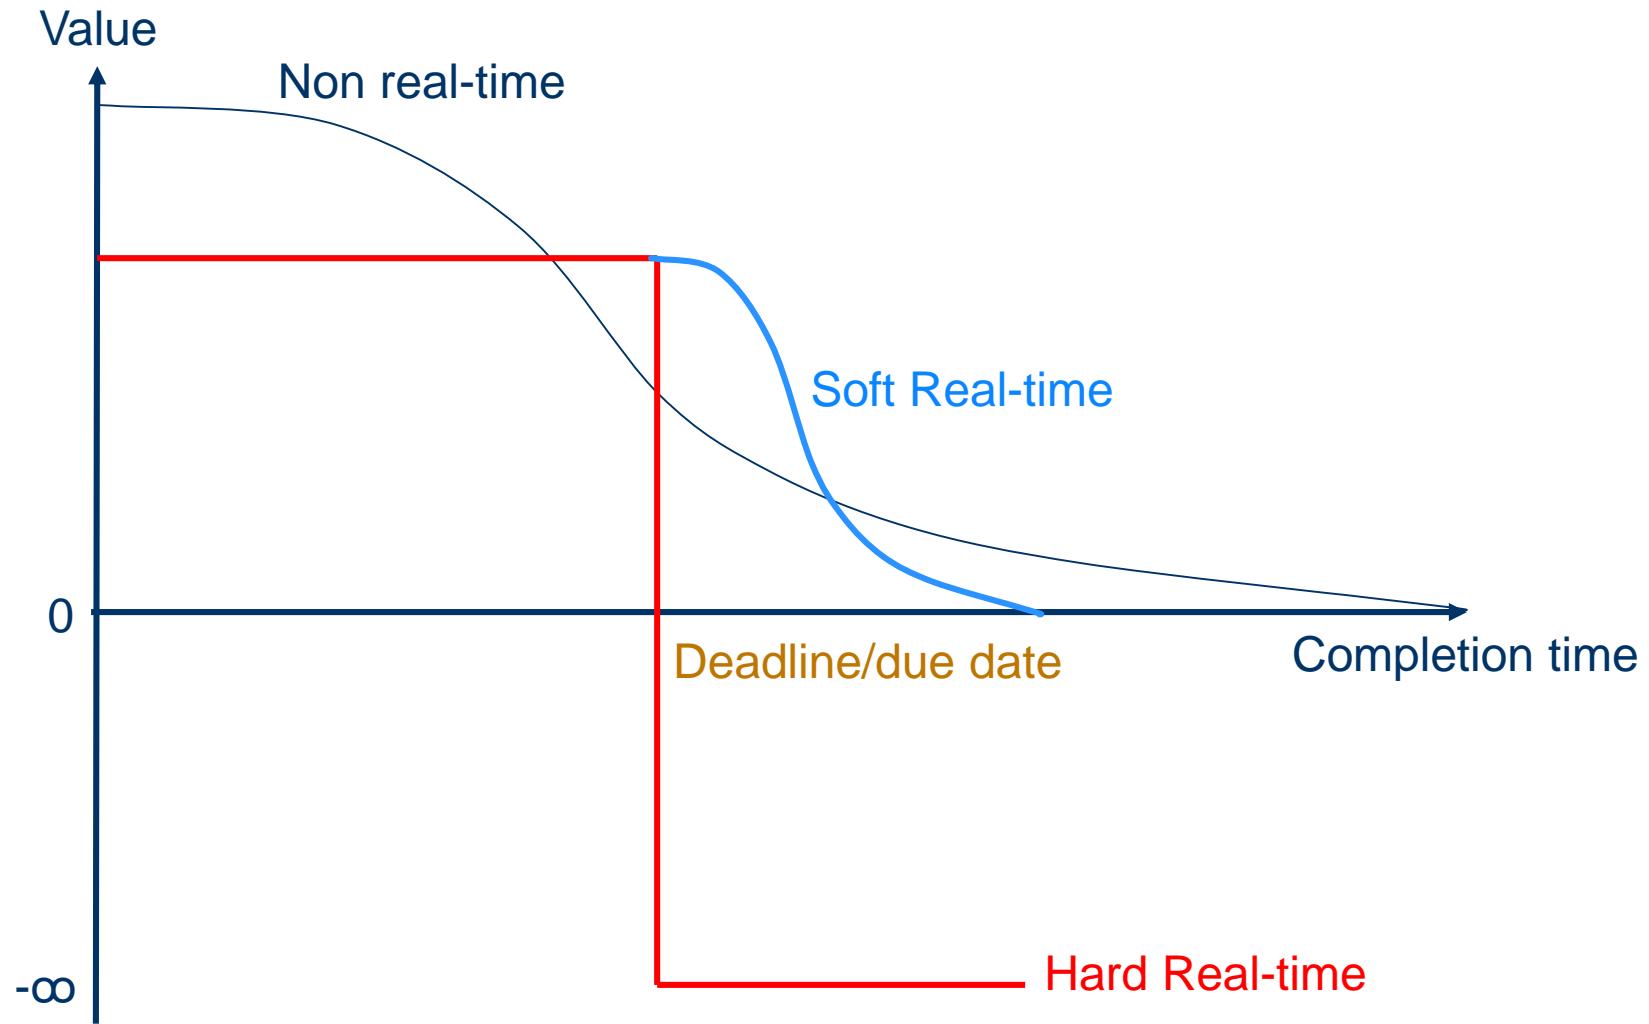
- Air Traffic Control Systems
- Networked Multimedia Systems
- Command Control Systems
- Space Shuttle
- Missile
- Train Control System
- Electronic Control System (Engine, Transmission, Headlight, ABS, etc)
- Weapons
- Factory Automation System

❖ Correctness of System

- Logically correct result of computation
- Time at which the results are produced



Real-time Systems



Real-Time Scheduling

	C_i	P_i
T_1	1	4
T_2	2	6

C_i : Computation Time
 P_i : Period



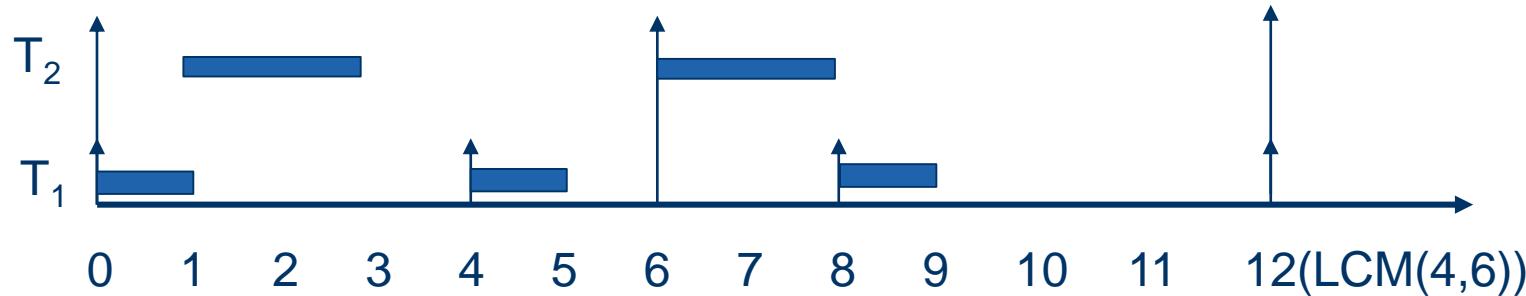
RM(Rate Monotonic) Scheduling

- Fixed priority
- Shorter period, higher priority



Not
Schedulable

Schedulable



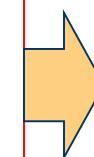
Real-Time Scheduling

	C_i	P_i
T_1	1	4
T_2	2	6
T_3	3	8



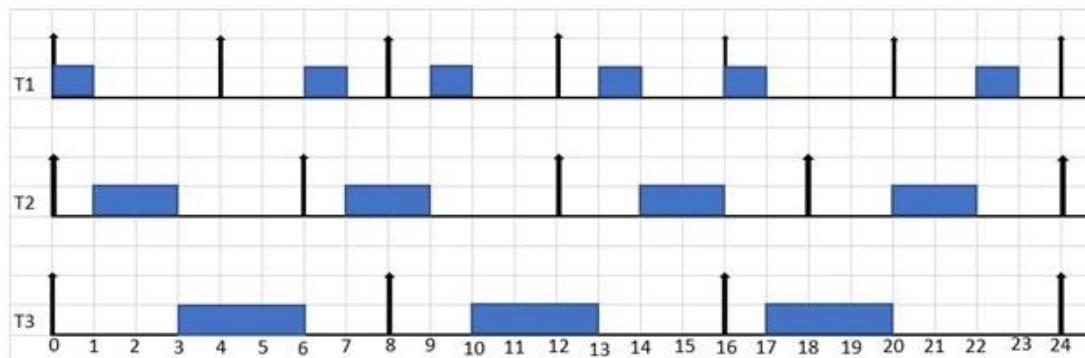
EDF(Earliest Deadline First) Scheduling

- Dynamic Priority
- Earlier Deadline, higher priority



Not
Scheduled

Schedulable



Schedulability Analysis

❖ Schedulability Analysis

- Total Utilization $U = \sum U_i = \sum (C_i/P_i)$
- If $(U \leq n(2^{1/n}-1))$, schedulable by RMS
- If $(U \leq 1)$, schedulable by EDF

❖ $\lim n(2^{1/n}-1) = \ln 2 \approx 0.693$